


Operating Systems

Part 8

Operating Systems

- Master controller for all of the activities that take place within a computer
- Basic Duties:
 - manage the physical resources of the system
 - load and execute programs
 - controlling I/O devices



5/8/2018 Sacramento State - Cook - CS&35 - Spring 2018 2

What is an operating system?

- The operating system is simply another program executing on the processor
- But, it runs in *privileged (or supervisor) mode*
 - knows about all the hardware in the computer
 - has it the ability to run special instructions
- Other programs run in *user mode*

5/8/2018 Sacramento State - Cook - CS&35 - Spring 2018 3

Interact with Applications

- The operating reserves many activities for itself – for stability and reliability
- These include:
 - input/output – keyboard, screen ports
 - memory
 - special registers
 - etc...

5/8/2018 Sacramento State - Cook - CS&35 - Spring 2018 4


Interact with Applications

- So, programs need to "talk" to the OS using *Application Program Interface (API)*
- Benefits:
 - makes applications faster and smaller
 - also makes the system more secure since apps do not directly talk to IO
 - Application → Operating System → IO

5/8/2018 Sacramento State - Cook - CS&35 - Spring 2018 5

Vector Tables

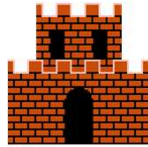
- Programs (and hardware) often need to talk to the operating system
- Examples:
 - software needs talk to the OS – in other words: API
 - USB port notifies the OS that a device was plugged in



5/8/2018 Sacramento State - Cook - CS&35 - Spring 2018 6

Vector Tables

- But how does this happen?
- The processor can be *interrupted* - alerted that something must be handled
- Each type interrupt has a unique number – which identifies the type of alert



5/8/2018

Sacramento State - Cook - CS&35 - Spring 2018

7

Vector Table



- When interrupted, the processor looks up the number in the *"vector table"*
- Table contains the address of the subroutine to execute
- The interrupt number is an index into this table

5/8/2018

Sacramento State - Cook - CS&35 - Spring 2018

8

How it Works

- When interrupted, the processor uses the interrupt number (index into the table) and looks up the address
- It then executes that address (*like a function you call in your Java programs*)



5/8/2018

Sacramento State - Cook - CS&35 - Spring 2018

9

How it Works

- These subroutines belong to the *kernal* – the core of the operating system
- So, software can interrupt itself with a specific number (*designated for software to use*) when it needs to talk to the operating system



5/8/2018

Sacramento State - Cook - CS&35 - Spring 2018

10

Instruction: Interrupt

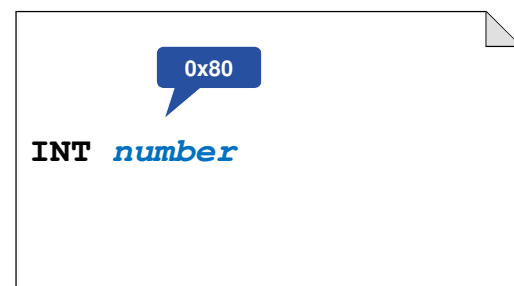
- Interrupt Instruction allows your program to "interrupt" itself and pass information to the operating system *kernal*
- How you use it
 1. fill registers with values that will tell Linux what to do
 2. call Linux by using interrupt 0x80 (or a special *software interrupt* instruction)

5/8/2018

Sacramento State - Cook - CS&35 - Spring 2018

11

Instruction: Interrupt (32-bit)



5/8/2018

Sacramento State - Cook - CS&35 - Spring 2018

12

Instruction: syscall (64-bit)

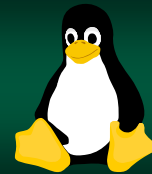
SYSCALL

Calls vector reserved for OS

5/9/2018

Sacramento State - Cook - CS&35 - Spring 2018

13

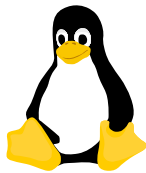


Linux System Calls

How software and hardware "talk"

Interrupts on the Linux

- Linux, like other operating systems communicate with applications using *interrupts*
- Applications do not know where (in memory) to contact the kernel – so they ask the processor to do it



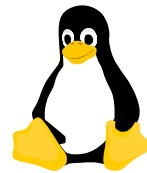
5/9/2018

Sacramento State - Cook - CS&35 - Spring 2018

15

How It Works

1. Fill the registers
2. Interrupt using 0x80 (or the special *software interrupt* instruction in 64-bit)
3. Any results will be stored in the registers



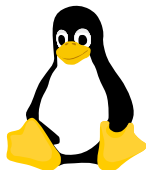
5/9/2018

Sacramento State - Cook - CS&35 - Spring 2018

16

Kernels are Simple!

- Linux only has **1** write and **1** read system call
- The location, number of bytes, and device only change
- It basically states "*write x many bytes from y to device z*"
- So, writing to the screen, a file, a port, etc...use the same call!



5/9/2018

Sacramento State - Cook - CS&35 - Spring 2018

17

How to Call Linux – 32 bit



- In the 32-bit version of Linux, the kernel uses eax to identify the system call
- Each call has a unique constant
- Interrupt 0x80** activates the kernel (well, the code that handles app requests)

5/9/2018

Sacramento State - Cook - CS&35 - Spring 2018

18

How to Call Linux – 32 bit



- Often a system call needs additional information to work
- Registers ebx, ecx, edx, etc... will contain this information

5/8/2018

Sacramento State - Cook - CS:35 - Spring 2018

19

Linux 32 – Sys Write

```

mov  $4, %eax
mov  $1, %ebx
mov  $address, %ecx
mov  $length, %edx
int  $0x80
    
```

Linux 32 call for WRITE

1 = Screen

How many bytes

Call Linux

5/8/2018

Sacramento State - Cook - CS:35 - Spring 2018

20

Linux 32: Sys Read

```

mov  $3, %eax
mov  $0, %ebx
mov  $address, %ecx
mov  $maxBytes, %edx
int  $0x80
    
```

Linux 32 call for READ

0 = Keyboard

Maximum number of bytes to read

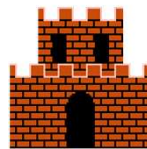
Call Linux

5/8/2018

Sacramento State - Cook - CS:35 - Spring 2018

21

How to Call Linux – 64 bit



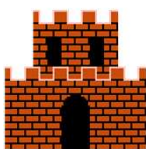
- The 64-bit version of Linux changed the system calls
- The rax register still holds the system call number
- However, the "kernel call" numbers are different now
- There are 329 total calls

5/8/2018

Sacramento State - Cook - CS:35 - Spring 2018

22

How to Call Linux – 64 bit



- Different registers are used to hold data
- The order is also quite different: *rdi, rsi, rdx, r10, r8*
- The new instruction *syscall* talks to the OS

5/8/2018

Sacramento State - Cook - CS:35 - Spring 2018

23

Some Linux 64 Calls

System Call	rax	rdi	rsi	rdx
read	0	fd (device)	address	max bytes
write	1	fd (device)	address	count
open	2	address	flags	mode
close	3	fd (device)		
get pid	39			
exit	60	error code		

5/8/2018

Sacramento State - Cook - CS:35 - Spring 2018

24

Linux 64: Sys Write

```
mov $1, %rax  
mov $1, %rdi  
mov $address, %rsi  
mov $length, %rdx  
syscall
```

Linux command
for WRITE

1 = Screen

Call Linux

5/9/2018

Sacramento State - Cook - CSC 35 - Spring 2018

25

Linux 64: Sys Read

```
mov $0, %rax  
mov $0, %rdi  
mov $address, %rsi  
mov $maxBytes, %rdx  
syscall
```

Linux command
for READ

0 = Keyboard

Maximum
number of bytes
to read

Call Linux

5/9/2018

Sacramento State - Cook - CSC 35 - Spring 2018

26