




Heaps & Priority Queues

Section 2.4




Heaps

Piles of data!

What is a heap?

- A *heap* is a binary tree, but a notable format to the nodes
- The value of a node is smaller (or larger) than **both** of its children
- Every subtree is a heap



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

3

Min and max-heaps


- A *min-heap*
 - stores smaller items (minimal items) at the top of the tree
 - larger items are stored at the bottom
- A *max-heap*
 - stores larger items (maximum items) at the top of the tree
 - smaller items are stored at the bottom

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

4


Min and max-heaps



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

5



Terminology Warning



- The Heap ADT is **not** the same as the operating system's heap
- ADT is a tree that stores "heavier" objects at the bottom
- The system heap is a complex interwoven linked list


6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

6

Heaps

- Heaps are *complete* binary trees
- Nodes are added in breadth-first order
- The resulting tree is always optimal and *balanced*



6/26/2019

Sacramento State - Summer 2019 - CS130 - Cook

7

Height of a Heap


- Let h be the height of the heap
- Let i be the depth of a node
- For all $i = 0, \dots, h - 1$
 - there are 2^i nodes of depth i
 - at depth $h - 1$, the internal nodes are to the left of the external nodes
 - The last node of a heap is the rightmost internal node of depth $h - 1$

6/26/2019

Sacramento State - Summer 2019 - CS130 - Cook

8

Height of a Heap



6/26/2019

Sacramento State - Summer 2019 - CS130 - Cook

9

Why Use a Heap?

- Heaps are not as fast as B-Trees, but are far simpler for maintaining balance than AVL Trees
- We will cover these soon...
- Filling logic makes them a good idea for arrays
 - any array can be turned into a heap
 - ... even if it *already* contains data
 - the data will need to be restructured to a proper heap
 - there is a sort algorithm based on this – *Heap Sort*

6/26/2019

Sacramento State - Summer 2019 - CS130 - Cook

10

Adding an Node


- Begin at next available position for a leaf
- Now the item needs to be *up-heaped*
 - move the entry up depending on its value until a correct position is found
 - as this is done, nodes are swapped entries from parent to child change position
 - since a heap *always* has height $O(\log n)$, upheap runs in $O(\log n)$ time

6/26/2019

Sacramento State - Summer 2019 - CS130 - Cook

11

Minheap: Adding a Node 13




6/26/2019

Sacramento State - Summer 2019 - CS130 - Cook

12

Upheaping 13




6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

13

Upheaping 13




6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

14

Upheaping 13: Done



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

15

Total Up-heap-val

- Just to make matters confusing, up-heap and down-heap are also known by various other terms – which are all valid
- These are some:
 - bubble-up
 - percolate-up
 - sift-up
 - heapify-up
 - cascade-up

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

16

Deleting a Node

- Deleting a node is quite different from adding
- Remember, heaps must maintain *completeness*
- So, the right-most leaf will be involved
- Deletion:
 - remove the node and replace it with the right-most leaf
 - now, this node needs to *down-heap* (moved down) to the correct location
 - since a heap *always* has height $O(\log n)$, down-heap runs in $O(\log n)$ time

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

17

Downheap Algorithm


- With a heap, every node has two children
 - as you downheap, you swap nodes
 - so which one do you select?
- Preserve the heap structure ← **vital**
 - on a min-heap, swap with the smallest child
 - on a max-heap, swap with the largest child

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

18

Minheap: Deleting 25




6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

19

Deleting 25: Replace




6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

20

Deleting 25: Downheaping




6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

21

Deleting 25: Downheaping




6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

22

Deleting 25: Downheaping



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

23

As You Expected...

- Just like up-heap, down-heap has several other, completely valid, names
- These are some:
 - bubble-down
 - percolate-down
 - sift-down
 - heapify-down
 - cascade-down

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

24


Merging Heaps



Whoa, this is easy!

Merging Heaps

- Merging two heaps is actually quite easy
- One approach is to read one heap into another....
 - read all the data from one heap and add it to the second
 - requires $O(n \log n)$



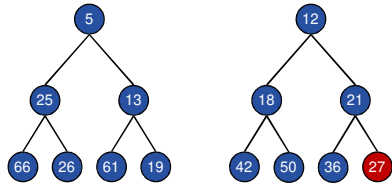
6/26/2019 Sacramento State - Summer 2019 - CS130 - Cook 26

Merging Heaps

- Or, we can create a new root
 - remember: every subtree in a heap – is a heap
 - so, both trees can be added as a left / right subtree
 - just grab a node at the base of one, make it the root, and downheap
 - requires $O(\log n)$

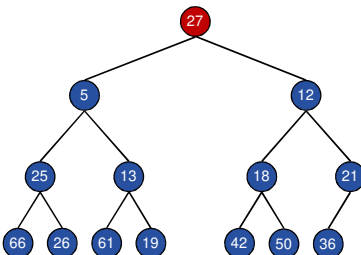
6/26/2019 Sacramento State - Summer 2019 - CS130 - Cook 27

Merge: Create New Root



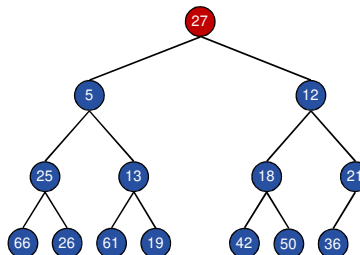
6/26/2019 Sacramento State - Summer 2019 - CS130 - Cook 28

Merge: Create New Root




6/26/2019 Sacramento State - Summer 2019 - CS130 - Cook 29

Merge: Downheap



6/26/2019 Sacramento State - Summer 2019 - CS130 - Cook 30

Merge: Downheap




6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

31

Merge: Downheap




6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

32


Merge: Done



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

33




Heaps and Arrays

Not a Heap O' Trouble

Heaps and Arrays

- Heaps are *complete*, balanced, binary trees
- This rigid, predictable, structure...
 - lends itself to being stored in an array
 - each node has a pre-ordained location



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

35

Heaps and Arrays

- Using an array, links between items...
 - are not explicitly stored
 - finding array index of an item's children and parent can be found using simple mathematics
- Heap ADTs only need to...
 - track the index of the end of the heap
 - all new items are added here – before upheap
 - and this is where the last item will be swapped for a deleted item (before it is downheaped)

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

36

Heap Array Math



Find	0 Indexed Array	1 Indexed Array
Parent of node i	$(i - 1) / 2$	$i / 2$
Left child of node i	$(2 * i) + 1$	$2 * i$
Right child of node i	$(2 * i) + 2$	$(2 * i) + 1$

6/26/2019

Sacramento State - Summer 2019 - CS130 - Cook

37

Heap in an Array





6/26/2019

Sacramento State - Summer 2019 - CS130 - Cook

38

Heap in an Array



6/26/2019

Sacramento State - Summer 2019 - CS130 - Cook

39



Queues can play favorites

Priority Queues

- A stack is first-in-last-out
- A queue is first-in-first-out
- A priority queue is *first-in-least-out*



6/26/2019

Sacramento State - Summer 2019 - CS130 - Cook

41

First-in-Least-Out

- "least" element is the first one removed
- If two items have the same "rank", items can be queued as normal
- The value that is used to determine an item's value is called a "key"



6/26/2019

Sacramento State - Summer 2019 - CS130 - Cook

42

What is the "Least" Item?

- Meaning of "least" is defined by the ADT
- It is an abstract term and does not mean "minimal"
 - so, "least" can be any way of ranking items
 - ...if the items are mathematically transitive
 - "least" can be the largest value
- Examples
 - by the smallest / largest value
 - size of the data (e.g. files)

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

43

Typical Priority Queue ADT

- Main Operations:
 - **add** (object) : add the object to the priority queue
 - **removeLeast** : removes and returns the least item
 - **getLeast** : returns the least element
 - **isEmpty** : returns true if the PQ is empty

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

44

Implementation

- Before we select a data structure to implement a priority queue, we should look how data will be used
- The goal is to get the **best time** efficiency with as little overhead
- Know what data will be stored will influence how the priority queue is implemented and, importantly, **how**



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

45

Implementation

- We can use a basic data structure
 - array
 - linked-list
 - tree
- Or another ADT
 - B-Tree
 - Queue
 - Heap

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

46

Implementation with an Array

- Unsorted array
 - adding requires $O(1)$
 - removing requires $O(n)$ – search and moving
- Sorted array
 - adding requires $O(n)$ – find location, move rest
 - removing requires $O(1)$ – if the head of the queue is at the array **end**
- Both approaches are inefficient

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

47

Implementation with a Linked List

- Unsorted linked list
 - adding requires $O(1)$
 - removing requires $O(n)$ – find and delete node
- Sorted linked list
 - adding requires $O(n)$ – find position and insert
 - removing requires $O(1)$ – just remove the head/tail
- Just as inefficient as pure arrays

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

48

Implementation with a Binary Tree

- Unbalanced B-Tree
 - adding requires $O(\log n)$ to $O(n)$
 - removing requires $O(\log n)$ to $O(n)$
- Balanced binary tree
 - adding requires $O(\log n)$
 - removing requires $O(\log n)$
 - rebalancing requires an extra $O(\log n)$ time

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

49

Implementation with a Heap

- A priority queue can be implemented as a heap
- Remember...
 - in a heap, all the items below a node have a greater value
 - if that value is used as a priority key, the items on the top of the heap is the top of the queue

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

50

Implementation with a Heap

- *Heaps naturally implement a priority queue*
- To enqueue an item...
 - just add it to the heap
 - it will up-heap to the correct position
 - requires $O(\log n)$
- To dequeue an item
 - just delete the root
 - requires $O(\log n)$ rebalance



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

51

Hybrid Implementations

- In some cases, the key value can have a minor range of values – possibly just a few
- Examples:
 - hospital triage – immediate, delayed, minor
 - computer processes – OS, application, GUI
- We can make clever hybrid structures that maximize efficiency

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

52

Hybrid Implementations

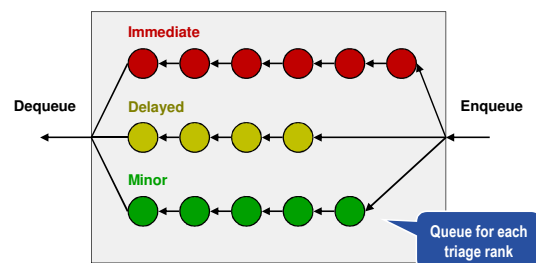
- If the key contains a small number of values, you can use multiple queues – one for each key value
- Basically, the priority queue, internally, will have an array of queues
- Adding/removing items will always be $O(1)$
 - $O(1)$ for the queue head
 - $O(1)$ for enqueue/dequeue (using a linked list)

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

53

Hospital Triage – Array of Queues



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

54

... But Heaps are Universal


- However, in most cases, the key values have **large** ranges
- For example, if the key is a 32-bit integer, do you want to create 4 million queues?
- Didn't think so....
- The pure Heap implementation works in all cases

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

55

Hospital Triage - Heap



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook


56

Heap Sort

Heap-a-rific algorithm!

Heap Sort

- *Heap Sort* is an ingenious sort algorithm that uses a max-heap to sort an array of elements
- It takes advantage of ...
 - heap is a natural priority queue
 - heaps **always** add / remove from the right-most leaf




6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

58

Heap Sort Works in Two Phases

- Phase 1: *Heapify*
 - the sort first converts the existing array into a heap
- Phase 2: Empty heap
 - removes all the nodes (treating it as priority queue)
 - sorted data is added to the end of the array



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

59

Implementation

- Both the "heap" and the remaining array can be used in memory at the same time
- The sorted array is stored at the empty space **after** the end of the heap
- This concept works for both Phase 1 and Phase 2

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

60

Phase 1: Array → Heap

- In Phase 1, we convert the array into a maxheap. This step is called *heapify*.
- Remember....
 - a heap can be stored in an array
 - so, we can just look at the array as a heap
 - ...but, its not quite a heap yet
 - data needs to be moved around until it is a heap

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

61

How do we convert it?

- First approach: *top-down*
 - start building the heap at the top of the array
 - iterate the variable *i* starting from the first element and build a heap above *i*
 - this is the easiest to conceptualize
- Second approach: *bottom-up*
 - fastest approach is to downheap all the leaves
 - leaves are always located at the second half of the array
 - so, we run the downheap, at the root, all the leaves

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

62

Phase 1: Heapify

```
n = count-1; //last item

while (n >= 0)
{
    heapDown(array, n, count-1)
    n--;
}
```

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

63

Phase 1: Heapify

```
heapify(array, count)
{
    last = count - 1;
    n = last; //last item

    while (n >= 0)
    {
        downHeap(array, n, last)
        n--;
    }
}
```

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

64

Phase 2: Root Deletion

- Now that the array is a *max*heap, the root contains the *maximum* item
- If we remove the root, we have the *last* item in a sorted array!
- OMG! Sooooo, awesome!



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

65

Phase 2: Root Deletion

- Remember, when we remove the root...
 - right-most leaf is moved to the root and downheaped into the correct position
 - this leaf position is now empty



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

66

Phase 2: Root Deletion

- We can put the root, we just removed, in this new empty space
- *What a sec!* We just put the largest item in the last position in the array
- The value, now at the root of the heap, is the second largest item in the array

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

67

Phase 2: Root Deletion

- So, to sort the array....
 - so, we just keep removing the root and placing it position where the leaf was located
 - the "heap" section of the array shrinks as the sorted array grows from the bottom
 - wow, that's easy!

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

68

Heap Sort Algorithm

```
last = count - 1;
heapify(array, 0, last);
while (last > 0)
{
    // swap root and array[last]
    downHeap(0, last - 1);
    last--;
}
```

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

69

Phase 1 – Top-down Convert

■ Heap
■ Array



Conceptual View


6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

70

Phase 1 – Top-down Convert

■ Heap
■ Array



Conceptual View



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook


71

Phase 1 – Top-down Convert

■ Heap
■ Array



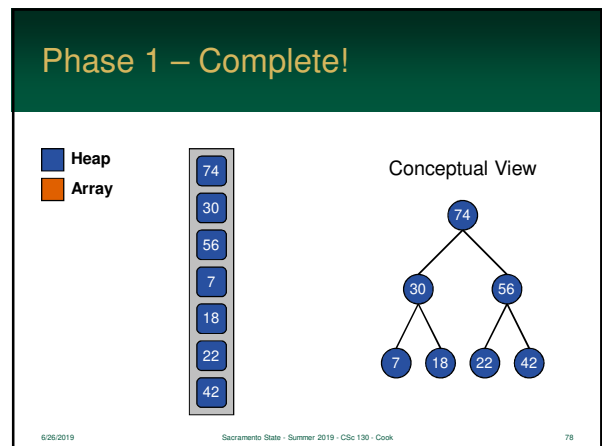
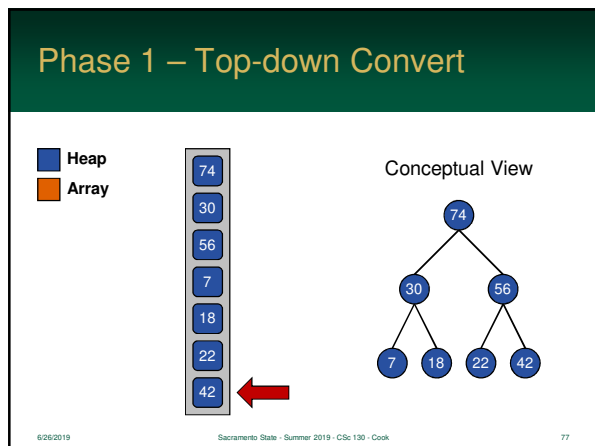
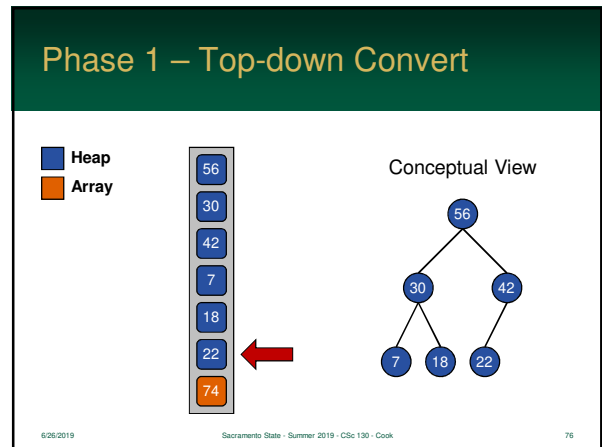
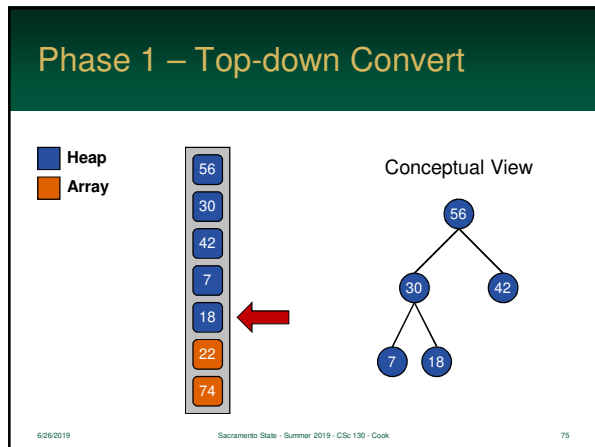
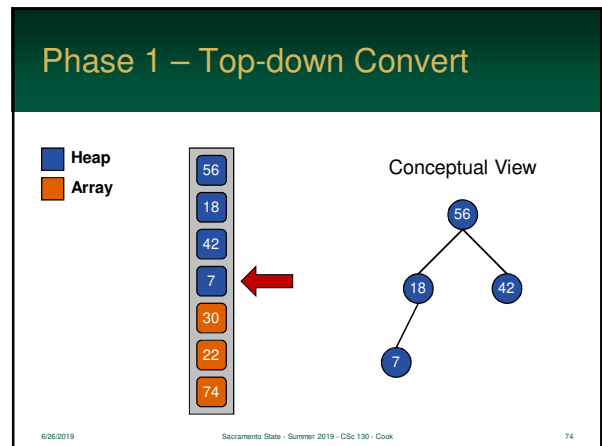
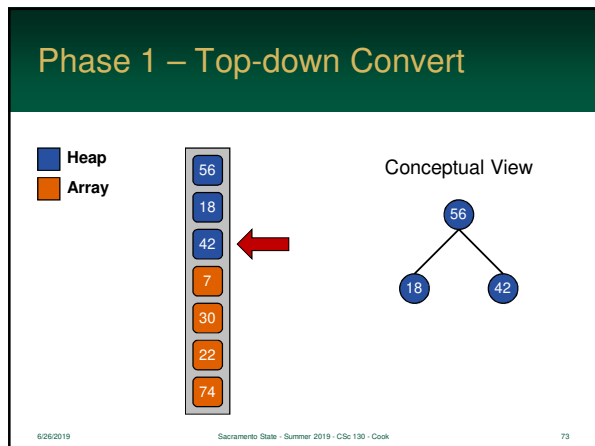
Conceptual View

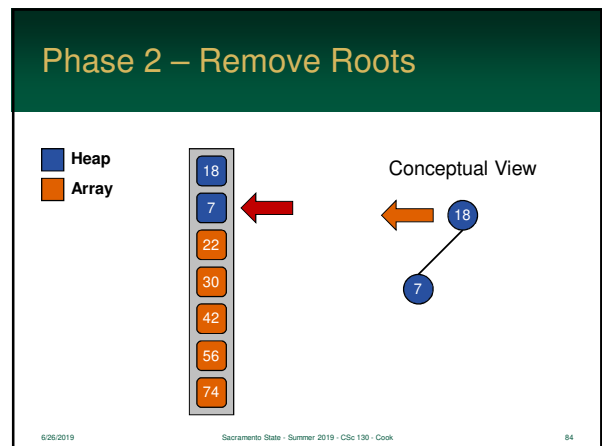
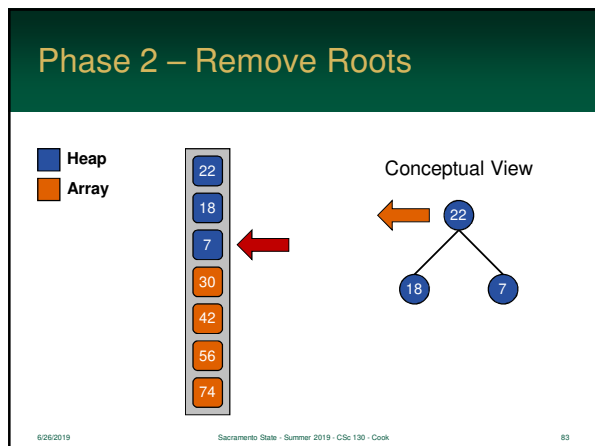
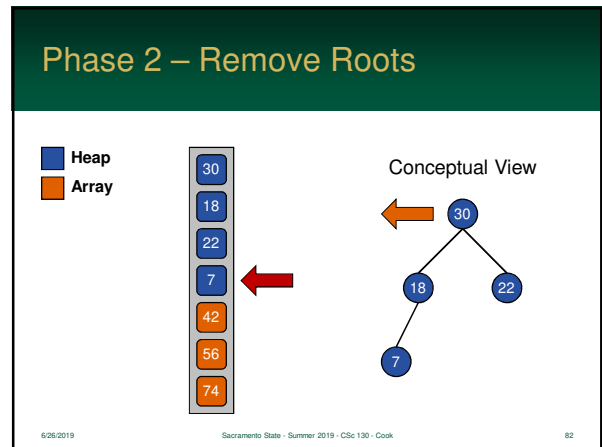
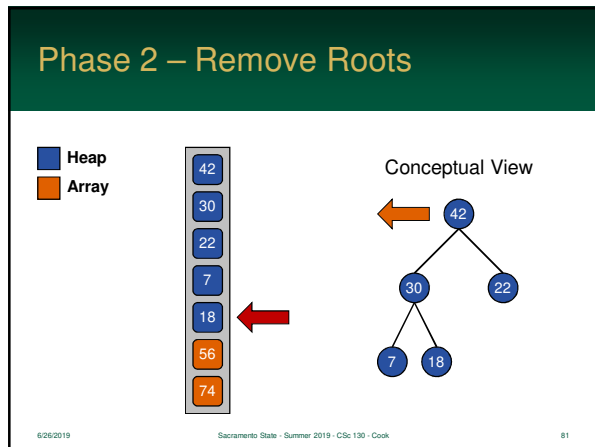
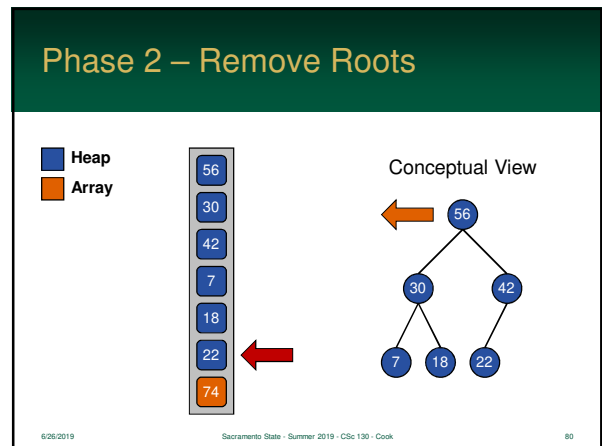
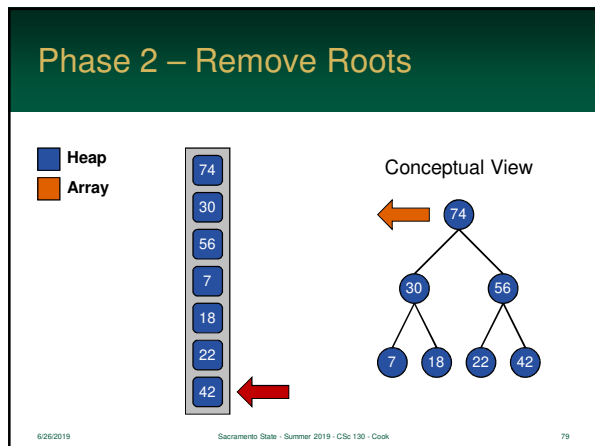


6/26/2019


Sacramento State - Summer 2019 - CSc 130 - Cook

72





Phase 2 – Remove Roots




6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

85

Phase 2 – Complete!



6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

86

Merge Sort vs. Heap Sort

- Heap-Sort allows us to sort any array in $O(n \log n)$ just like Merge-Sort & Quicksort
- However, there is no overhead
 - Heap-Sort can be sorted in-place, meaning auxiliary storage is $O(1)$
 - Merge-Sort, however, requires $O(n)$
 - Quick-Sort can become $O(n^2)$

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

87

Merge Sort vs. Heap Sort

- However, in some cases, the recursive nature of Merge Sort is better
 - easy to distribute to multiple computers
 - Heap-Sort uses the entire array – not online
- But...in the Real World, it gets complex
 - you can cut an array into sub-lists, send them to different machines which Heap-Sort them
 - ... and then you Merge

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

88

Heap Sort Summary

Heap Sort	
Time Average	$O(n \log n)$
Time Best	$O(n \log n)$
Time Worst	$O(n \log n)$
Auxiliary space	$O(1)$
Stable	No – Equal element order not preserved
Online?	No

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

89

Summary of Sorting Algorithms

Algorithm	Best	Average	Worst	Aux. Storage
Bubble Sort	$O(n^2)$	$O(n^2)$	$O(n^2)$	$O(1)$
Selection Sort	$O(n^2)$	$O(n^2)$	$O(n^2)$	$O(1)$
Insertion Sort	$O(n)$	$O(n^2)$	$O(n^2)$	$O(1)$
Shell Sort	$O(n \log n)$	$O(n^{5/4})$	$O(n^{3/2})$	$O(1)$
Merge Sort	$O(n \log n)$	$O(n \log n)$	$O(n \log n)$	$O(n)$
Quick Sort	$O(n \log n)$	$O(n \log n)$	$O(n^2)$	$O(1)$
Heap Sort	$O(n \log n)$	$O(n \log n)$	$O(n \log n)$	$O(1)$

6/26/2019

Sacramento State - Summer 2019 - CSc 130 - Cook

90